

# Course Overview

ITM424:  
Computer Systems  
1<sup>st</sup> Lecture, Sep. 7, 2023

**Instructors:**  
Moon Hea-Eun (Ocean)

# Overview

- Introductions
- Big Picture
  - Course theme
  - Five realities
  - How the course fits into the CS/ECE/INI curriculum
- Logistics

# Instructor



- ocean.moon@seoultech.ac.kr
- mayfly74@gmail.com

## Bio

- 2014 ~ Now : NSHC Inc. Red Alert Research Institute Director
- 2002 ~ 2014 : NetMan Co., Ltd. Research Institute Director

## Latest Activities

- 2020.02 : IAEA Conference Speaker in Austria
- 2019.08 : DEFCON ICS/SCADA CTF Organizer
- 2019.08 : CISS SCADA CTF – Winner
- 2018.08 : DEFCON ICS/SCADA CTF Organizer
- 2018.07 : IOT Singapore CTF – 2<sup>nd</sup> Winner
- 2017.11 : CODEBLUE Hack2Win 2017 – Winner
- 2017.08 : CISS SCADA CTF – Winner
- 2015.11 : CODEBLUE Hack2Win 2015 – Winner
- 2015.10 : Whitehat Contest Final
- 2014.10 : Whitehat Contest Final 3<sup>rd</sup>
- 2013.08 : DEF CON, Capture The Flag 8<sup>th</sup>
- 2010.10 : ISEC 2010 Capture The Flag winner
- 2010.08 : DEF CON, Capture The Flag 6<sup>th</sup>

# Instructor - Research

- Offensive Research
- Bug hunting research
- More Focus on IoT, ICS/SCADA
- Found over 255 real world 0-day vulnerabilities
- Specialize in Red Teaming Exercise
- Provide Consulting, Test Bed, CTF, Penetration Test, Vulnerability Assessment and Training Service

# The Big Picture

# Course Theme:

## (Systems) Knowledge is Power!

### ■ Systems Knowledge

- How hardware (processors, memories, disk drives, network infrastructure) plus software (operating systems, compilers, libraries, network protocols) combine to support the execution of application programs
- How you as a programmer / hacker can best use these resources

### ■ Useful outcomes from taking this class

- Become more effective programmers
  - Able to find and eliminate bugs efficiently
  - Able to understand and tune for program performance
- Prepare for later “systems” classes in CS & ECE
  - Compilers, Operating Systems, Networks, Computer Architecture, Embedded Systems, Storage Systems, etc.

# It's Important to Understand How Things Work

- **Why do I need to know this stuff?**

- Abstraction is good, but don't forget reality

- **Most CS and CE courses emphasize abstraction; so do we!**

- Abstract data types
- Asymptotic analysis

- **But it's helpful to understand what abstractions build upon and their limits!**


- Especially in the presence of bugs
- Need to understand details of underlying implementations
- Sometimes the abstract interfaces don't provide the level of control or performance you need

# Great Reality #1:

## Ints are not Integers, Floats are not Reals

### ■ Example 1: Is $x^2 \geq 0$ ?

- Float's: Yes!



- Int's:
  - $40000 * 40000 \rightarrow 1600000000$
  - $50000 * 50000 \rightarrow ?$

### ■ Example 2: Is $(x + y) + z = x + (y + z)$ ?

- Unsigned & Signed Int's: Yes!
- Float's:
  - $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
  - $1e20 + (-1e20 + 3.14) \rightarrow ??$



# Great Reality #2:

## You've Got to Know Assembly

- **Chances are, you'll never write programs in assembly**
  - Compilers are much better & more patient than you are
- **But: Understanding assembly is key to machine-level execution model**
  - Behavior of programs in presence of bugs
    - High-level language models break down
  - Tuning program performance
    - Understand optimizations done / not done by the compiler
    - Understanding sources of program inefficiency
  - Implementing system software
    - Compiler has machine code as target
    - Operating systems must manage process state
  - Creating / fighting malware
    - x86 assembly is the language of choice!

# Great Reality #3: Memory Matters

## Random Access Memory Is an Unphysical Abstraction

### ■ Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated

### ■ Memory referencing bugs especially pernicious

- Effects are distant in both time and space

### ■ Memory performance is not uniform

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

# Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

```
fun(0) --> 3.14  
fun(1) --> 3.14  
fun(2) --> 3.13999998664856  
fun(3) --> 2.000000061035156  
fun(4) --> 3.14  
fun(6) --> Segmentation fault
```

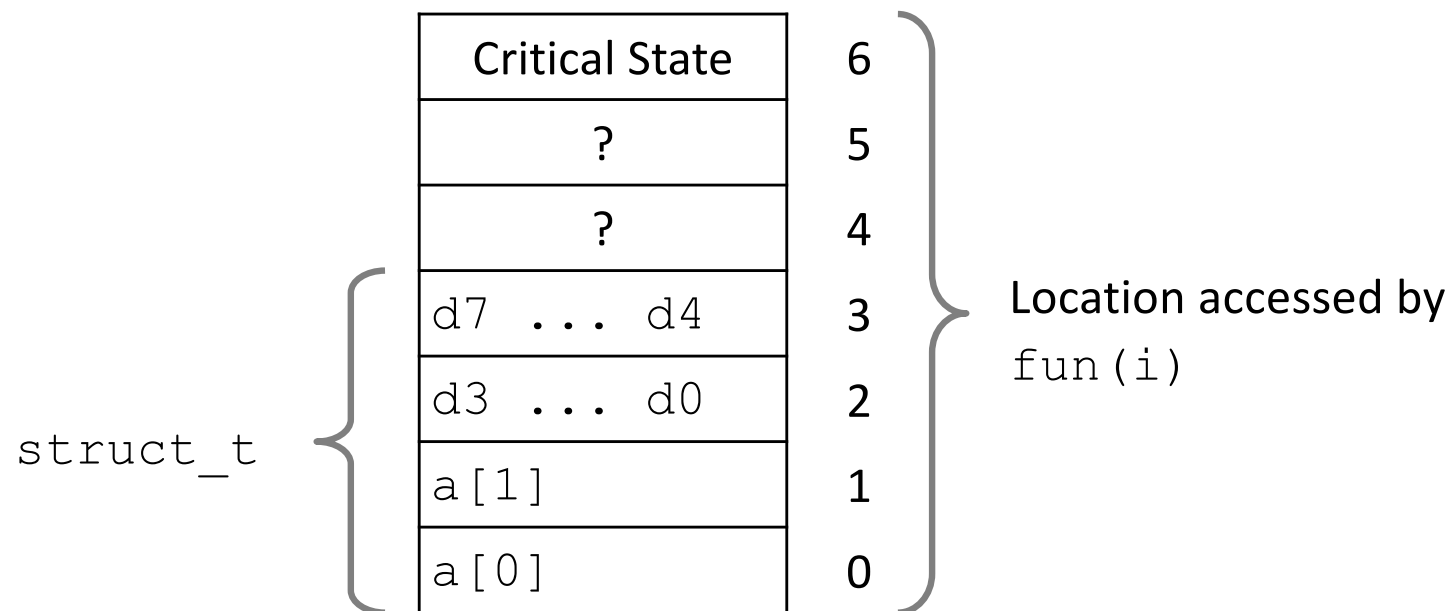
- Result is system specific

# Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;
```

```
fun(0) --> 3.14
fun(1) --> 3.14
fun(2) --> 3.1399998664856
fun(3) --> 2.00000061035156
fun(4) --> 3.14
fun(6) --> Segmentation fault
```

Explanation:



# Memory Referencing Errors

## ■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

## ■ Can lead to nasty bugs

- Whether or not bug has any effect depends on system and compiler
- Action at a distance
  - Corrupted object logically unrelated to one being accessed
  - Effect of bug may be first observed long after it is generated

## ■ How can I deal with this?

- Program in Java, Ruby, Python, ML, ...
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. Valgrind)

# Great Reality #4: There's more to performance than asymptotic complexity

- **Constant factors matter too!**
- **And even exact op count does not predict performance**
  - Easily see 10:1 performance range depending on how code written
  - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- **Must understand system to optimize performance**
  - How programs compiled and executed
  - How to measure program performance and identify bottlenecks
  - How to improve performance without destroying code modularity and generality

# Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

4.3ms

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

81.8ms

2.0 GHz Intel Core i7 Haswell

- Hierarchical memory organization
- Performance depends on access patterns
  - Including how step through multi-dimensional array

# Logistics



# Primary Textbook

- Randal E. Bryant and David R. O'Hallaron,
  - *Computer Systems: A Programmer's Perspective*, **Third Edition** (CS:APP3e), Pearson, 2016
  - <https://csapp.cs.cmu.edu>
  - This book really matters for the course!
    - How to solve labs
    - Practice problems typical of exam problems
  - Electronic editions available (*Don't get paperback/international version!*)
- Note: All textbooks have errors
  - Don't panic if you see something that seems wrong
  - Come talk to us about it if you can't make it make sense

# Recommended reading

- Brian Kernighan and Dennis Ritchie,
  - *The C Programming Language*, Second Edition, Prentice Hall, 1988
  - Everyone calls this book “K&R”
  - Guide to C by the designers of the language
  - Well-written, concise
  - A little dated
    - Doesn't cover additions to C since 1988 (that's thirty years ago...)
    - Casual about issues we consider serious problems now

# If you want more books about C

- C for Programmers with an introduction to C11
  - Paul and Harvey Deitel
  - Opposite of K&R: modern, verbose
  - Lots of worked-out examples
  - Ugly code style (compare readability to K&R)
- 21<sup>st</sup> Century C
  - Ben Klemens
  - Supplement to full C textbooks: goes into the corners of the language
  - Opinionated
  - First half is about how to *build* C programs in the Unix environment
    - So, if you want to understand the Makefiles we give you...
- Learn C the Hard Way
  - Zed A. Shaw
  - Extremely opinionated
  - Also has lots of worked-out examples
  - Only book I can find that takes “undefined behavior” seriously enough

# Course Components

## ■ Lectures

- Higher level concepts
- Provide in-depth understanding of an aspect of systems
- Programming and measurement

## ■ Assignments

- Reinforce concepts
- 2023-2\_ITM424\_Assignment.doc

## ■ Exam

- Test your understanding of concepts & mathematical principles
- Covers content from the whole semester

## ■ Grading system

- Homework (30, Mid-term (30) and Final (40)

# Programs and Data

## ■ Topics

- Bit operations, arithmetic, assembly language programs
- Representation of C control and data structures
- Includes aspects of architecture and compilers

# The Memory Hierarchy

## ■ Topics

- Memory technology, memory hierarchy, caches, disks, locality
- Includes aspects of architecture and OS

# Exceptional Control Flow

## ■ Topics

- Hardware exceptions, processes, process control, Unix signals, nonlocal jumps
- Includes aspects of compilers, OS, and architecture

# Networking, and Concurrency

## ■ Topics

- High level and low-level I/O, network programming
- Internet services, Web servers
- concurrency, concurrent server design, threads
- I/O multiplexing with select
- Includes aspects of networking, OS, and architecture



# Schedule

## ■ Topics

W	Date	Contents
1	2023.09.08	Linux Overview
2	2023.09.15	C Programming Language
3	2023.09.22	Bits, Bytes, & Integers
4	2023.09.29	Machine Level Programming (Basic, Control)
5	2023.10.06	Machine Level Programming (Procedures, Data)
6	2023.10.13	Machine Level Programming (Advanced)
7	2023.10.20	Mid Term Exam
8	2023.10.27	Design and Debugging
9	2023.11.03	The Memory Hierarchy
10	2023.11.10	Linking
11	2023.11.17	Processes and Multitasking, Exceptional Control Flow
12	2023.11.24	System Level I/O
13	2023.12.01	Network Programming
14	2023.12.08	Network Programming
15	2023.12.15	Final Exam

# Getting Help

## ■ Class Web page:

- e-Class
- Complete schedule of lectures, exams, and assignments
- Copies of lectures, assignments, exams, solutions
- FAQ

## ■ TA

- Info.

## ■ Email

- Feel free to ask any questions

Welcome  
and Enjoy!